

# Lab: Graph Neural Networks (GNN)

Instructor: Fragkiskos Malliaros  
TA: Vahan Martirosyan

January 22, 2026

## Description

In this lab, you will learn to use a powerful representation learning family: *Graph Neural Networks*. In particular, you will implement a Graph Convolution Network (GCN) [1] to solve both a node classification task and a graph classification task..

## Description

**Graph Neural Networks (GNNs)** are powerful deep learning architectures for representation learning of graph data. They achieve state-of-the-art results in a wide variety of tasks by deriving very informative node embeddings that incorporate both graph structure and node feature information. This is due to their specific functioning, expressed as a recursive message passing scheme, where they encode information from nodes and pass it along the edges of the graph.

For the specific case of **Graph Convolution Networks**: we consider a graph  $\mathcal{G}$  with feature matrix  $\mathbf{X}$ , adjacency matrix  $\mathbf{A}$ , and diagonal degree matrix  $\mathbf{D}$ . Let  $N$  be the number of nodes,  $C$  the number of input features,  $F$  of output features, and  $\mathbf{Z}$  the output. Note that,  $\mathbf{Z} \in \mathbb{R}^{N \times F}$ ,  $\mathbf{X} \in \mathbb{R}^{N \times C}$ ,  $\mathbf{W} \in \mathbb{R}^{C \times F}$ , where  $\mathbf{W}$  is a weight matrix.

The output of one GCN layer is obtained as follows:

$$H^{(l+1)} = f(H^{(l)}, A) = \sigma(\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} H^{(l)} \mathbf{W}^{(l)})$$

with  $H_0 = \mathbf{X}$ ,  $H_L = \mathbf{Z}$  and  $\sigma$  often chosen to be the ReLU function. Also,  $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$  and  $\tilde{\mathbf{D}} = \mathbf{D} + \mathbf{I}$  (adding self loops). Each element in  $H_{(l+1)}$  can thus be written as  $h_i^{(l+1)} = \sigma(\sum_{j \in N(i)} \frac{1}{c_{ij}} W^{(l)} h_j^{(l)})$ .

## Part I: Node Classification

The dataset, Amazon Photo, is a segment of the Amazon co-purchase graph. Nodes represent goods and edges indicate that two goods are frequently bought together. Node features are bag-of-words encoded product reviews, and class labels are given by the product category.

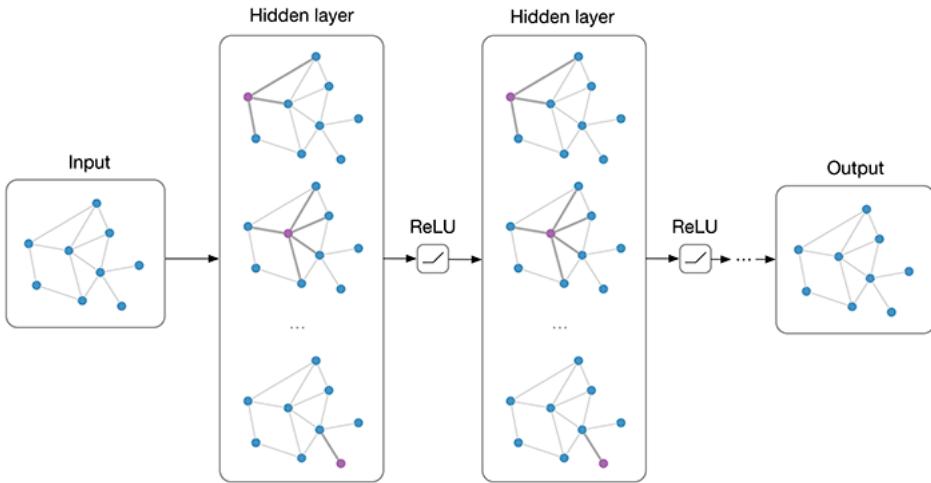


Figure 1: Multi-layer Graph Convolutional Network (GCN)

**Exercise 1: Pre-processing**

Before you start, make sure you have the required packages (e.g. DGL<sup>1</sup>, torch, sklearn) installed.

1. Make sure the dataset is correctly imported in the main function. You will be using the popular Deep Graph Library (DGL) to implement a GNN architecture. DGL makes it easy to implement GNNs since it provides many state-of-the-art GNN layers and modules. You should retrieve different key properties: the feature matrix  $X$ , the label vector  $y$ , the number of classes, the number of features and the number of nodes. Careful, the graph is not a networkx instance and commands are slightly different for DGL graphs...
2. Use the `split_dataset()` method to divide the graph into training/validation/test sets for node classification.

Now that the data is pre-processed, let's define, train and evaluate a GNN model on it.

**Exercise 2: Learning**

1. Construct a Graph Convolution Network (GCN) inside the `GNN_model` class. You can follow the architecture in figure 1. You do not have to redefine it from scratch. GCN blocks are already defined inside the `dgl` function `GraphConv`. *Note:* We use the pytorch backend by default but `dgl` also operates with tensorflow.
2. Define an optimizer, a loss function and train the GCN model inside `train()`.
3. Although it is a great approach, GCN has some potential limits. Could you think of some? For instance, what could happen if we stack a lot of GCN layers? You can use the interactive visualization of GCN available at: <https://distill.pub/2021/understanding-gnns/#interactive>

**Part II: Graph Classification**

In this section, we will follow a similar idea but for a graph classification task. Instead of having a single graph where you want to assign a label to nodes, you have a set of graphs, and would like to assign

<sup>1</sup><https://docs.dgl.ai/#>

a label to each one of them. For instance, in biology, we have the so-called protein-protein interaction networks. Each protein of is represented as a graph, and we would like to infer the function of this protein.

### Exercise 3

1. Load the following dataset `dgl.data.TUDataset(name='ENZYMES')`. Create a training, validation and test set (Careful the dataset format is particular).  
You can use `dgl.data.utils.split_dataset` to split the dataset.  
Finally, use the `GraphDataLoader` function to batch graphs.
2. Create a GNN model to perform the graph classification task. Reminder: we want to derive a graph representation, not just node representation.
3. Train this new model and evaluate it, similarly to part I.

## References

- [1] Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*, 2016.